

Ganesh Navre

• +91 9579225968 • ganeshnavre1@gmail.com • Pune • [Github](#) • [LinkedIn](#)

SUMMARY

AI Engineer with hands-on experience in Generative AI, LLMs, and end-to-end ML pipelines, building production-ready AI applications. Skilled in **Python, HuggingFace Transformers, LangChain, RAG, prompt engineering, FastAPI, and model deployment**. Experienced in model training, fine-tuning, recommendation systems, and NLP. Built an AI-powered chatbot using Streamlit and Transformers (distilgpt2) for context-aware personalized recommendations, boosting engagement and order value. Experienced in **collaborating on AI projects, documenting experiments, and translating business problems into intelligent solutions** in fast-paced environments.

TECHNICAL SKILLS

AI / LLMs & Generative AI

- **HuggingFace Transformers** (distilgpt2, BERT, Sentence Transformers) — text generation & NLP
- **LangChain** — Chains, Agents, Memory, Tool use, Prompt templates
- **OpenAI GPT API, Claude API** — Prompt engineering, few-shot, chain-of-thought techniques
- **RAG (Retrieval-Augmented Generation)** — FAISS, semantic search, embedding pipelines
- Fine-tuning basics, context-aware generation, personalized recommendation systems

Machine Learning & MLOps

- **Scikit-learn, TensorFlow, PyTorch (basics)** — Regression, Classification, Random Forest, Feature Engineering
- **Model Evaluation** (Accuracy, Precision, Recall, F1) — model versioning & performance monitoring
- **FastAPI, Streamlit** — REST API development & ML model deployment
- **Docker (basics), GitHub Actions** — CI/CD for ML pipelines & automated workflows

Programming & Tools

- **Python (Advanced)** — Pandas, NumPy, Matplotlib, Seaborn
- **SQL** (Joins, Subqueries, Window Functions) — MySQL, SQL Server
- Git, GitHub, VS Code, Jupyter Notebooks, Google Colab

Cloud & Vector Databases

- **AWS S3 / SageMaker (basics), Azure ML Studio** — Cloud AI/ML platforms (learning & hands-on)
- **FAISS (basics)** — Vector search, semantic similarity, embedding-based retrieval

Data Analysis & Visualization

- **Power BI, Tableau** — Dashboards, KPIs, Visual Storytelling
- **Excel/Google Sheets** — Pivot Tables, Advanced Formulas
- EDA, KPI Tracking, Trend & Anomaly Detection, Dashboard Reporting & Automation

Statistics & Experimentation

- Descriptive Statistics, Correlation Analysis, Probability Basics
- **A/B Testing fundamentals** (hypothesis, control vs variant, metrics comparison)
- Experiment result analysis & interpretation, KPI monitoring and anomaly detection

EXPERIENCE

AI / Data Annotator & Analytics Content Reviewer (Freelance)

Innodata Inc. | Remote

Duration: Nov 2025 – Present

- Reviewed and evaluated LLM-generated outputs for accuracy, safety, and alignment, contributing directly to **AI model quality improvement** and training data pipelines.
- Analyzed annotation quality metrics and trends to **identify recurring errors, process gaps, and improvement opportunities**.
- Cleaned, validated, and reviewed **large-scale structured and unstructured datasets** used for model training and fine-tuning.
- Applied **analytical judgment** to classify AI edge cases under strict compliance and privacy guidelines.
- Documented findings and provided structured feedback to improve AI data process efficiency.

Data Science Intern

MentorMind | Uber x upGrad

Duration: Jul 2025 – Dec 2025

- Built an **end-to-end ML regression pipeline** using historical Uber trip data to predict ride fares.
- Performed **data cleaning, preprocessing, and EDA** to uncover key pricing drivers and improve model performance.
- Supported experiment analysis, distinguishing **data-related vs model-related metric deviations**, supporting ML debugging workflows.
- Documented assumptions, experiments, and findings; delivered final model, Jupyter notebooks, and project report.

KEY PROJECTS

AI-Powered Swiggy Chatbot – Smart Meal Assistant

Tech Stack: Python, Streamlit, HuggingFace Transformers (distilgpt2), LangChain, Generative AI, Pandas, Recommendation Systems

- Developed a Generative AI chatbot using **distilgpt2 and HuggingFace Transformers** to provide personalized meal recommendations based on user taste, order history, mood, time, and weather context.
- Implemented **filters, trending dishes, chef specials, combo suggestions, and an interactive cart**, enhancing user engagement and driving potential increase in average order value and retention.
- Engineered a **context-aware prompt pipeline** for real-time LLM inference on CPU with minimal latency, demonstrating practical MLOps and deployment skills.

TripPulse – Ride Delay Prediction & Data Analysis

Tech Stack: Python, Pandas, NumPy, Matplotlib, Seaborn, Scikit-learn, Random Forest

- Performed end-to-end **ML pipeline on 1500+ ride records**: EDA, feature engineering, model training, and evaluation.
- Engineered time-based and route-level features; trained **Random Forest model achieving 89% accuracy**.
- Monitored model performance metrics and validated outputs to detect anomalies and ensure consistency across data splits.

Commodity Price Analysis

Tech Stack: SQL (MySQL, SQL Server), Relational Database Design, Window Functions

- Designed and optimized a relational database integrating commodity, region, and price data.
- Evaluated price variations and regional trends across India (2019–2021) to identify cost-sensitive commodities and guide strategic decisions.

EDUCATION

Dr. Babasaheb Ambedkar Marathwada University | Chhatrapati Sambhajanagar, IN

Bachelor of Science | Jun 2022 – Jul 2025

CERTIFICATION

- Certified in Data Science & Artificial Intelligence (AI/ML) – UpGrad, 2025
- SQL (Basic) – HackerRank, May 2025
- Data Analytics Job Simulation – Deloitte (Forage), Mar 2025
- Business Analysis Foundations – Microsoft & LinkedIn Learning, Apr 2025

ADDITIONAL INFORMATION

Languages

- English (Proficient)
- Hindi (Native)
- Marathi (Native)